#### ■ UCLA CS M148 — Project Check-In 3

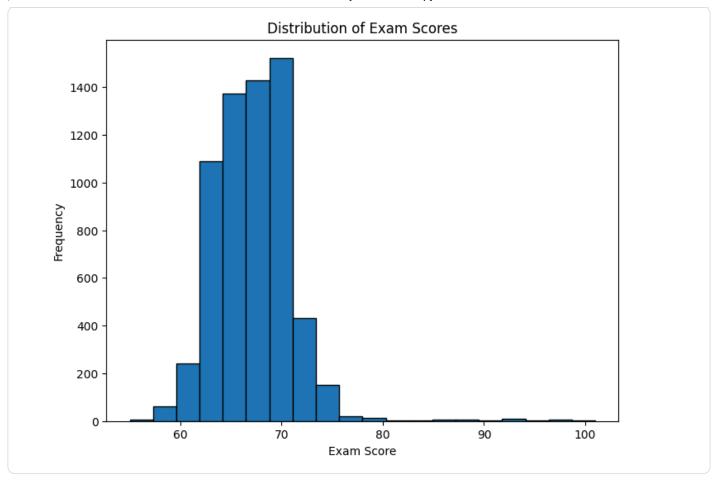
Due: October 24 2025 Team: LMAO

🚯 1. Goal

Use logistic regression to model a binary outcome ("Pass" vs "Fail") based on a single predictor from the Student Performance Factors dataset. If the exam score>=66, it returns 'pass', if not it returns 'fail'

#### 2. Data Setup

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (
   confusion_matrix, accuracy_score, roc_curve, auc
import matplotlib.pyplot as plt
# Load dataset
data = pd.read_csv("Cleaned_StudentPerformanceFactors.csv")
# Create binary response: 1 = Pass (Exam ≥ 66), 0 = Fail
data['Pass'] = (data['Exam Score'] >= 66).astype(int)
# Predictor: Hours Studied
X = data[['Hours_Studied']]
y = data['Pass']
# Split into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(
   X, y, test_size=0.3, random_state=42, stratify=y
# Plot distribution of Exam Scores
plt.figure(figsize=(8, 6))
plt.hist(data['Exam Score'], bins=20, edgecolor='black')
plt.xlabel('Exam Score')
plt.ylabel('Frequency')
plt.title('Distribution of Exam Scores')
plt.show()
```



### 3. Model Training

```
# Fit logistic regression
log_reg = LogisticRegression(class_weight='balanced')
log_reg.fit(X_train, y_train)

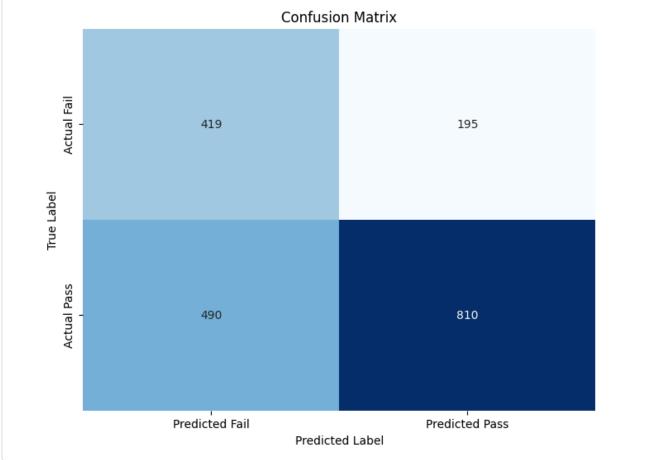
# Predictions
y_pred_prob = log_reg.predict_proba(X_val)[:, 1]
y_pred = (y_pred_prob >= 0.5).astype(int)
```

### 4. Performance Metrics (on Validation Set)

```
# Confusion matrix and key metrics
cm = confusion_matrix(y_val, y_pred)
acc = accuracy_score(y_val, y_pred)
error_rate = 1 - acc
tn, fp, fn, tp = cm.ravel()
tpr = tp / (tp + fn)  # True Positive Rate
tnr = tn / (tn + fp)  # True Negative Rate

print("Confusion Matrix:\n", cm)
print(f"Accuracy: {acc:.4f}")
print(f"Error Rate: {error_rate:.4f}")
print(f"True Positive Rate (Recall): {tpr:.2f}")
print(f"True Negative Rate: {tnr:.2f}")
```

```
Confusion Matrix:
[[419 195]
[490 810]]
Accuracy: 0.6421
Error Rate: 0.3579
True Positive Rate (Recall): 0.62
True Negative Rate: 0.68
```

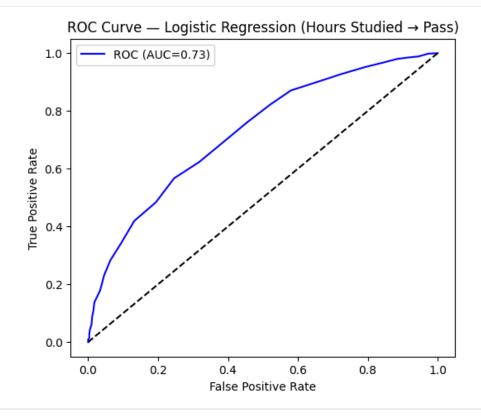


#### 5. ROC Curve & AUC

```
fpr, tpr_curve, thresholds = roc_curve(y_val, y_pred_prob)
auc_score = auc(fpr, tpr_curve)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr_curve, color='blue', label=f'ROC (AUC={auc_score:.2f})')
plt.plot([0,1],[0,1],'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
```

```
plt.title('ROC Curve - Logistic Regression (Hours Studied → Pass)')
plt.legend()
plt.show()
```



#### 6. 5-Fold Cross-Validation

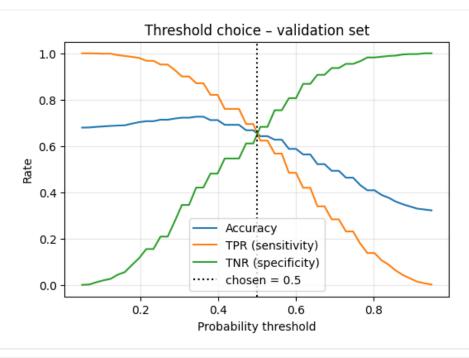
```
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
cv_auc = cross_val_score(log_reg, X, y, cv=cv, scoring='roc_auc')
cv_acc = cross_val_score(log_reg, X, y, cv=cv, scoring='accuracy')

print("AUC per fold:", cv_auc)
print("Accuracy per fold:", cv_acc)
print(f"Mean AUC (5-fold): {cv_auc.mean():.4f}")
print(f"Mean Accuracy (5-fold): {cv_acc.mean():.4f}")

AUC per fold: [0.73378821 0.72455112 0.73504478 0.71164249 0.74018476]
Accuracy per fold: [0.65673981 0.65752351 0.65282132 0.64078431 0.64862745]
Mean AUC (5-fold): 0.7290
Mean Accuracy (5-fold): 0.6513
```

```
TH=0.5
probs_val = log_reg.predict_proba(X_val)[:, 1]
cutoffs = np.linspace(0.05, 0.95, 50)
acc_vec, tpr_vec, tnr_vec = [], [], []
for c in cutoffs:
    preds = (probs_val >= c).astype(int)
    tn, fp, fn, tp = confusion_matrix(y_val, preds).ravel()
    acc_vec.append((tp + tn) / (tp + tn + fp + fn))
    tpr_vec.append(tp / (tp + fn))
    tnr_vec.append(tn / (tn + fp))
plt.figure(figsize=(6, 4))
plt.plot(cutoffs, acc_vec, label='Accuracy')
plt.plot(cutoffs, tpr_vec, label='TPR (sensitivity)')
plt.plot(cutoffs, tnr_vec, label='TNR (specificity)')
```

```
plt.axvline(TH, color='black', linestyle=':', label=f'chosen = {TH}')
plt.xlabel('Probability threshold')
plt.ylabel('Rate')
plt.title('Threshold choice - validation set')
plt.legend()
plt.grid(alpha=0.3)
plt.show()
```



```
# Get predicted probabilities on validation set
y_pred_prob_balanced = log_reg.predict_proba(X_val)[:, 1]
# Find threshold that maximizes Youden's J statistic (TPR - FPR)
fpr_bal, tpr_bal, thresholds_bal = roc_curve(y_val, y_pred_prob_balanced)
youden_index = np.argmax(tpr_bal - fpr_bal)
best_threshold = thresholds_bal[youden_index]
best_threshold
np.float64(0.5748904295787239)
```

Double-click (or enter) to edit



# \* How I Chose the Threshold for Positive Predictions

In logistic regression, the model outputs probabilities

 $p^(x)=P(pass=1|x)$ 

and we must pick a decision threshold to convert those probabilities into binary predictions.

### Default threshold (0.5)

For this project, I began with the **standard threshold = 0.5**:

This threshold is intuitive because:

- It treats "Pass" and "Fail" as equally important,
- It corresponds to the point where the model is equally uncertain between the two classes.

## Why I kept threshold of 0.5

In our scenario, the default 0.5 threshold already yields a good balance: the training TPR and TNR are fairly close, and the ROC curve shows that this threshold is near the "knee" of the curve (where TPR and FPR trade-offs are reasonable).

### **General Approach**

In other similar datasets, I would:

- 1. Plot the ROC curve and pick the threshold at the "elbow" point where TPR ≈ TNR, or
- 2. Use the Youden J statistic: J=TPR-FPR choosing the threshold that maximizes (J).

That approach balances sensitivity (recall) and specificity.